

Essential Steps for CNN and their needed implementations for RISC V:

1. Input Layer

- This is where the image data goes in.
- Each image is represented as a grid of pixel values.
- Unsure what size image we want, 28x28 is the typical but it might need to be reduced for computation time

RISC V Implementation:

- Load image data from memory using LW (load word) or LB (load byte) if stored as bytes.
- Store it into registers or arrays in memory for processing.
- Looping is needed for multi pixel images.

2. Convolutional Layer

- This layer applies filters (kernels) to scan parts of the image.
- Each filter is a small matrix that slides across the image matrix.
- The filter detects patterns like edges, shapes, and textures depending on values inside filter matrix.
- The result is a new version of the image called a feature map, highlighting the details that each filter corresponds to.

RISC V Implementation:

- Nested loops to slide the filter across the image.
- Perform multiply-accumulate (MAC) operations using MUL and ADD.
- Put values into memory.

3. Activation Function

- After convolution, relu activation function is applied.
- Keeps positive values and sets negatives to zero.
- Basically needed (boosted accuracy from 12% to ~50% when softmax wasnt used).

RISC V Implementation:

- Check register and if < 0 , it should be forced to zero.

4. Flattening

- The feature maps are still in a grid format (like an image).
- Flatten them into a single long vector to pass into the next layers.

RISC V Implementation:

- Loop the 2D array and append these values into a new 1D array.

5. Fully Connected Layer (Dense Layer)

- This layer connects every neuron to every other neuron.
- It takes the flattened values and learns to recognize high-level patterns (like digits, faces, or objects).
- It uses weights and biases to adjust importance for different features.

RISC V Implementation:

- Load the weights and inputs, then begin summations.
- Loops for all values.

6. Output Layer

- This is the layer that makes the prediction, last layer must represent the number of outputs (MNIST digits needs 10 neurons on last layer).
- Apply softmax or sigmoid depending on what is easier to implement (both are above 90% accuracy).

RISC V Implementation:

- Approximate e for softmax/sigmoid usage
- Output is a logit until softmax is applied (softmax makes it a probability)
- 10 value array for storing logits
- Usage of argmax is needed to get value and index (index is the digit its making the prediction for)